



EU AI Act Compliance Copilot

Link to the Chatbot: <https://aegis-eu-chatbot.vercel.app/>

THE CHALLENGE

AI regulation is now a product-management concern as much as a legal one. The EU AI Act takes a risk-based approach: what a system must do to comply depends on its purpose, its users, the people it affects, and its place in the AI value chain. Product teams commonly face three questions:

- **Does this apply to us?** Teams struggle to tell whether a new or existing AI feature falls into a regulated category.
- **What exactly must we do?** Turning legal text into concrete product, engineering, documentation and monitoring actions is hard.
- **Are we still compliant?** Obligations apply in phases and the dates have shifted. Prohibitions and GPAI duties already apply, transparency rules largely apply from August 2026, and the 2026 AI Omnibus moved standalone high-risk obligations to December 2027.

OUR SOLUTION

Aegis, the EU AI Act Compliance Copilot, is a conversational assistant that helps teams assess, prepare for and manage compliance for every AI product. Instead of reading hundreds of pages of regulation, a product manager simply describes the system in plain language, for example: *"We use an AI model to screen job applications and rank candidates."* The Copilot then guides them through a structured assessment.

Describe → Classify → Identify obligations → Assess gaps → Action plan

Classify Determines whether a system is likely to be prohibited, high-risk, subject to transparency rules, GPAI-related, or minimal risk.	Map obligations Translates legal requirements into practical duties across risk management, data governance, documentation, oversight and security.
Check compliance Asks targeted questions about the product, data, users and processes to uncover missing controls.	Recommend actions Produces a prioritized checklist, for example: add human oversight, document training data, enable logging, disclose AI use to users.
Generate evidence Drafts risk-assessment checklists, governance registers and questions for legal and security teams.	Stay current Tracks changing implementation dates and Commission guidance so products can be reassessed over time.

WHO IT SERVES

Product managers build compliance into design rather than adding it at launch. **Engineering and data teams** turn requirements into technical controls. **Legal and compliance teams** speed up first-level reviews and standardize assessments. **Business leaders** understand regulatory exposure and prioritize remediation.

THE VALUE WE DELIVER

- **Less effort:** automates the initial research and assessment that normally needs significant cross-functional work.
- **Earlier risk detection:** moves AI governance into product development, before gaps appear after deployment.
- **Accessible regulation:** converts legal language into simple, role-specific questions and actions.
- **Consistency at scale:** offers a repeatable assessment across an organization's entire AI portfolio.

From "Is our AI compliant?" to "Here is exactly what we need to do next."

OUR CORE PRINCIPLE

The Copilot does not replace lawyers or regulators. It is a first-line governance partner that helps teams understand the Act, identify likely obligations, surface gaps and prepare the right questions and evidence for formal review.

APPENDIX

Rules applied by the rules engine

Legal basis: Regulation (EU) 2024/1689 (the AI Act) as amended by Regulation (EU) 2026/1744 (the Digital Omnibus on AI, cited below as “Omnibus”). Rule pack version 2026.10.08. This appendix summarises the rule base; the complete rule base (64 classification rules, 50 obligations, 33 controls, every condition and its Official Journal page reference) is in the project repository.

How the engine decides

- **Rules in three-valued logic:** Rules combine facts with AND, OR and NOT. A rule that depends on an unknown fact is reported as “possible” and the missing fact becomes a question. Screening facts that are almost always “no”, such as prohibited practices, are assumed “no” and shown to the user as assumptions.
- **The Act's own order:** Scope, roles, prohibited practices, high-risk, role shifts (Article 25), transparency (Article 50), general-purpose AI models, then obligations, dates and fines.
- **Deterministic and traceable:** The same facts and date always give the same result, and every result carries its article and Official Journal page.

Classification rules

Scope (6 rules)

Rule	Test	Law
SCOPE_NOT_AI	Not an AI system	Art. 3(1)
SCOPE_NO_EU	No EU connection	Art. 2(1)
SCOPE_MILITARY	Military, defence or national security only	Art. 2(3)
SCOPE_RESEARCH	Scientific research only	Art. 2(6)
SCOPE_PREMARKET	Pre-market development only	Art. 2(8)
SCOPE_PERSONAL	Personal, non-professional use	Art. 2(10)

Roles (6 rules)

Rule	Test	Law
ROLE_PROVIDER	Provider	Art. 3(3)
ROLE_DEPLOYER	Deployer	Art. 3(4)
ROLE_IMPORTER	Importer	Art. 3(6)
ROLE_DISTRIBUTOR	Distributor	Art. 3(7)
ROLE_MANUFACTURER	Product manufacturer	Art. 2(1)(e)
ROLE_GPAI_PROVIDER	General-purpose AI model provider	Art. 3(63)

Prohibited practices (10 rules; the last two were added by the Omnibus and apply from 2 December 2026)

Rule	Test	Law
P5_A_MANIPULATION	Subliminal or manipulative techniques	Art. 5(1)(a)
P5_B_VULNERABILITIES	Exploiting vulnerabilities	Art. 5(1)(b)
P5_C_SOCIAL_SCORING	Social scoring	Art. 5(1)(c)
P5_D_CRIME_PROFILING	Predicting crime from profiling alone	Art. 5(1)(d)
P5_E_FACE_SCRAPING	Untargeted facial-image scraping	Art. 5(1)(e)
P5_F_EMOTIONS_WORK_SCHOOL	Emotion inference at work or in education	Art. 5(1)(f)
P5_G_BIOMETRIC_CATEGORISATION	Biometric categorisation of sensitive attributes	Art. 5(1)(g)
P5_H_REALTIME_RBI	Real-time remote biometric identification for law enforcement	Art. 5(1)(h)
P5_BA_INTIMATE_IMAGERY	Non-consensual intimate imagery	Omnibus Art. 5(1)(ba), (1a), (1b)
P5_BB_CSAM	Child sexual abuse material	Omnibus Art. 5(1)(bb), (1a)

High-risk (30 rules: 25 Annex III rules and 5 core rules)

Annex III has one rule per use case: the three biometric points (1(a) remote biometric identification, 1(b) biometric categorisation, 1(c) emotion recognition) and the 22 use cases of areas 2 to 8: 2 critical infrastructure; 3(a) education admission; 3(b) education evaluation; 3(c) education level; 3(d) education proctoring; 4(a) recruitment; 4(b) work

management; 5(a) public benefits; 5(b) credit scoring; 5(c) life health insurance; 5(d) emergency triage; 6(a) law enforcement victim risk; 6(b) law enforcement polygraph; 6(c) law enforcement evidence; 6(d) law enforcement offending risk; 6(e) law enforcement profiling; 7(a) migration polygraph; 7(b) migration risk; 7(c) migration applications; 7(d) migration identification; 8(a) justice; 8(b) elections.

The core rules choose between the Annex III and Annex I routes and apply the Article 6(3) exception. `d.annex3Candidate` is true when any Annex III rule fires; `d.exceptionHolds` is the result of the exception rule.

Rule	Condition	Law
HR_EXCEPTION_6_3	<code>d.annex3Candidate AND NOT profilesPeople AND (narrowProceduralTask OR improvesCompletedHumanWork OR detectsDecisionPatternsOnly OR preparatoryTaskOnly) AND NOT materiallyInfluencesDecisions</code>	Art. 6(3)-(4)
HR_PROFILING_OVERRIDE	<code>d.annex3Candidate AND profilesPeople</code>	Art. 6(3), last subparagraph
HR_ANNEX3	<code>d.annex3Candidate AND NOT d.exceptionHolds</code>	Art. 6(2); Annex III
HR_ANNEX1_SECTION_A	<code>safetyComponentOrProduct AND annex1Legislation is a law listed in Annex I, Section A (toys, lifts, radio equipment, medical devices and others) AND thirdPartyConformityAssessmentRequired AND NOT onlyNonSafetyFunctions</code>	Art. 6(1); Omnibus Art. 6(1a)-(1c); Annex I, Section A
HR_ANNEX1_SECTION_B	<code>safetyComponentOrProduct AND annex1Legislation is a law listed in Annex I, Section B (machinery, motor vehicles, civil aviation, marine equipment and others) AND thirdPartyConformityAssessmentRequired AND NOT onlyNonSafetyFunctions</code>	Art. 6(1); Omnibus Art. 2(2); Omnibus Annex I, Section B point 21

Role shifts (4 rules)

Rule	Test	Law
ROLE_ART25_REBRAND	Rebranding makes you the provider	Art. 25(1)(a)
ROLE_ART25_MODIFY	Substantial modification makes you the provider	Art. 25(1)(b)
ROLE_ART25_PURPOSE	Changing the purpose makes you the provider	Art. 25(1)(c)
ROLE_ART25_MANUFACTURER	Product manufacturer is the provider	Art. 25(3)

Transparency (5 rules)

Rule	Test	Law
T50_1_INTERACTION	Tell people they are dealing with an AI	Art. 50(1)
T50_2_MARKING	Mark generated content as AI-generated	Art. 50(2); Omnibus Art. 111(4)
T50_3_EMOTION_BIOMETRIC	Inform people exposed to emotion recognition or biometric categorisation	Art. 50(3)
T50_4A_DEEPFAKE	Disclose deep fakes	Art. 50(4)
T50_4B_PUBLIC_INTEREST_TEXT	Disclose AI-generated public-interest text	Art. 50(4)

General-purpose AI models (3 rules)

Rule	Test	Law
GPAI_PROVIDER	General-purpose AI model provider	Art. 3(63); Art. 53
GPAI_SYSTEMIC_RISK	General-purpose AI model with systemic risk	Art. 51; Art. 52; Art. 55
GPAI_OPEN_SOURCE_EXEMPT	Open-source exemption for part of Article 53	Art. 53(2); Art. 54(6)

Application dates

What	Original	Now applies
Chapters I and II: AI literacy and prohibited practices	2 February 2025	2 February 2025
New prohibitions (intimate imagery, child abuse material)	n/a	2 December 2026
General-purpose AI model obligations (Chapter V)	2 August 2025	2 August 2025
Penalties (Chapter XII, except Article 101)	2 August 2025	2 August 2025
Transparency obligations (Article 50)	2 August 2026	2 August 2026
Commission fines on general-purpose AI model providers (Article 101)	2 August 2026	2 August 2026
High-risk obligations for Annex III systems (Chapter III, Sections 1-3)	2 August 2026	2 December 2027
High-risk obligations for Annex I products (Chapter III, Sections 1-3)	2 August 2027	2 August 2028

Transitional rules: Article 111(2) (a high-risk system already on the market is caught only after significant design changes; systems for public authorities by 2 August 2030), Article 111(3) (general-purpose AI models placed on the market before 2 August 2025: 2 August 2027) and Article 111(4) (Article 50(2) marking for generative systems already on the market: 2 December 2026).

Obligations (50) by group

Group	Duties
Basics	AI literacy
Prohibited practices	Stop or redesign: prohibited practice; Stop or redesign: newly prohibited practice (from 2 December 2026)
High-risk: provider	Risk management system; Data and data governance; Technical documentation; Record-keeping (logging); Transparency and instructions for use; Human oversight by design; Accuracy, robustness and cybersecurity; Quality management system; Keep documentation and logs; Conformity assessment before launch; EU declaration of conformity and CE marking; Register in the EU database; Name, contact and accessibility; Post-market monitoring; Serious incident reporting; Corrective actions and cooperation with authorities; Appoint an EU authorised representative; Document the 'not high-risk' assessment; Register the system even though it is not high-risk
High-risk: deployer	Use as instructed; Assign competent human oversight; Relevant and representative input data; Monitor operation and report risks and incidents; Keep the logs; Inform workers before workplace use; Register use in the EU database (public authorities); Data protection impact assessment; Tell people they are subject to the system; Fundamental rights impact assessment; Right to an explanation
High-risk: importers and distributors	Importer checks before placing on the market; Distributor checks before making available
Value chain	You have become the provider; Written agreements with suppliers
Transparency	Tell users they are interacting with an AI; Mark generated content as AI-generated; Inform people exposed to emotion recognition or biometric categorisation; Disclose deep fakes; Disclose AI-generated public-interest text
General-purpose AI models	Model technical documentation; Information for downstream providers; Copyright policy; Public summary of training content; Appoint an EU authorised representative; Notify the Commission (systemic risk); Model evaluation and systemic-risk mitigation; Incident reporting and cybersecurity for the model

Fine tiers

- Up to EUR 35 million or 7% of total worldwide annual turnover, whichever is higher (Art. 99(3)).
- Up to EUR 15 million or 3% of total worldwide annual turnover, whichever is higher (Art. 99(4); Omnibus Art. 99(4)(da)).
- Up to EUR 15 million or 3% of total worldwide annual turnover, whichever is higher, imposed by the Commission from 2 August 2026 (Art. 101).
- Up to EUR 7.5 million or 1% of total worldwide annual turnover, for incorrect or misleading information to authorities (Art. 99(5)).
- Set by each Member State's penalty rules, with no fixed ceiling in the Act (Art. 99(1); Omnibus Art. 99(1)).
- For SMEs and small mid-caps each fine is capped at the lower of the percentage and the amount.

Confidence, open questions and gap check

- **Confidence:** High if the classification cannot change however the unknown facts turn out; Medium if two outcomes remain possible; Low if more, or if fewer than three facts are known. A judgement call, such as relying on the Article 6(3) exception, caps it at Medium.
- **Open questions:** The unknown facts that would change the classification, duties or dates, most influential first; whether it is an AI system and whether it has an EU connection come first.
- **Gap check:** 33 control questions, each evidencing one or more obligations. Missing or partial controls become actions: P0 (duty already in force and control missing, or a prohibited practice), P1 (duty in force with a partial control, or a missing control for a duty starting within 18 months), P2 (otherwise).
- **What-if:** A second set of facts is assessed and compared: tiers and duties added or removed, and dates moved.